

CS 570
Fall 2008
HW4 Solution (revised)

1) Suppose you are a consultant for the networking company Clunet, and they have the following problem. The network that they are currently working on is modeled by a connected graph $G=(V, E)$ with n nodes. Each edge e is a fiber-optic cable that is owned by one of two companies- creatively named X and Y - and leased to Clunet. Their plan is to choose a spanning tree T of G , and upgrade the links corresponding to the edges of T . Their business relations people have already concluded an agreement with companies X and Y stipulating a number k so that in the tree T that is chosen, k of the edges will be owned by X and $n-k-1$ of the edges will be owned by Y . Clunet management now faces the following problem: It is not at all clear to them whether there even exists a spanning tree meeting these conditions, and how to find one if it exists. So this is the problem they put to you: give a polynomial time algorithm that takes G , with each edges labeled X or Y , and either (i) returns a spanning tree with exactly k edges labeled X , or (ii) reports correctly that no such tree exists.

The algorithm:

Line 1	Remove all edges labeled X from G to get a new graph G' , which has k' disconnected components.
Line 2	If $k' > k + 1$
Line 3	return "no such tree exists"
Line 4	Else
Line 5	add exactly $k'-1$ edges labeled X in G' to make G' connected
Line 6	remove all edges labeled Y in G'
Line 7	add exactly $k-k'+1$ edges labeled X in G' subject to the constraint that no circle is formed in G' . If cannot find such edges, return "no such tree exists".
Line 8	add exactly $n-k-1$ edges labeled Y in G' subject to the constraint that no circle is formed in G'
Line 9	return G' , which is a spanning tree with exactly k edges labeled X .

2) Let us say that a graph $G = (V, E)$ is a near-tree if it is connected and has at most $n + 8$ edges, where $n = |V|$. Give an algorithm with running time $O(n)$ that takes a near-tree G with costs on its edges, and returns the minimum spanning tree of G . You may assume that all edge costs are distinct.

Algorithm:

```

Line 1:      While (G is not a tree)
Line 2:          find a circle in G
Line 3:          find the largest edge  $e$  in that circle
Line 4:          delete that edge  $e$  from G

```

3) A group of network designers at the communications company CluNet find themselves facing the following problem. They have a graph $G = (V, E)$, in which the nodes represent sites that what to communicate. Each edge e is a communication link, with a given available bandwidth b_e .

For each pair of nodes $u, v \in V$, they want to select a single u - v path P on which this pair will communicate. The *bottleneck rate* $b(P)$ of the path P is the minimum bandwidth of any edge it contains; that is, $b(P) = \min_{e \in P} b_e$. The *best achievable bottleneck rate* for the pair u, v in G is simply the maximum, over all u - v paths P in G , of the value $b(P)$.

It's getting to be very complicated to keep track of a path for each pair of nodes, and so one of the network designers makes a bold suggestion: Maybe one can find a spanning tree T of G so that for every pair of nodes u, v , the unique u - v path in the tree actually attains the best achievable bottleneck rate for u, v in G . (In other words, even if you could choose any u - v path in the whole graph, you couldn't do better than the u - v path in T .) This idea is roundly heckled in the offices of CluNet for a few days, and there's a natural reason for the skepticism: each pair of nodes might want a very different-looking path to maximize its bottleneck rate; why should there be a single tree that simultaneously makes everybody happy? But after some failed attempts to rule out the idea, people begin to suspect it could be possible.

Show that such a tree exists, and give an efficient algorithm to find one. That is: give an algorithm constructing a spanning tree T in which, for each $u, v \in V$, the bottleneck rate of the u - v path in T is equal to the best achievable bottleneck rate for the pair u, v in G .

Answer: Use a modified Prim's Algorithm in which you grow the spanned tree (until complete) with a node v that maximizes the attachment cost.

1. Initially $S = \{s\}$ and $d(s) = 0$
2. While $S \neq V$

3. Select a node v not in S with at least one edge from S for which $d'(v) = \max(l_e)$ for $e(u,v): u \in S$
4. Add v to S
5. End while

Starting from any node s would result in the same tree being spanned, so spanning tree T always contains the best achievable bottleneck rate paths for u to v in G .

Let the Case $|S| = 1; S = \{s\}$ and $d(s) = 0$

Assume it holds when $|S| = k$ for some value of $k \geq 1$. We now grow S to size $k+1$ by adding the node v . Let (u, v) be the final edge on our s - v path. Call this path P_v . Suppose we have another path to v that is shorter. It would have crossed from S to $V-S$ at some point, where the boundary edge would be from u' to v' . However, based on our greedy step in the algorithm, the shorter path u to v was picked instead, so it is contradiction to have a shorter path going through u' and v' to v .

4) In trying to understand the combinatorial structure of spanning trees, we can consider the space of all possible spanning trees of a given graph, and study the properties of this space. This is a strategy that has been applied to many similar problems as well.

Here is one way to do this. Let G be a connected graph, and T and T' two different spanning trees of G . We say that T and T' are *neighbors* if T contains exactly one edge that is not in T' , and T' contains exactly one edge that is not in T .

Now, from any graph G , we can build a (large) graph H as follows. The nodes of H are the spanning trees of G , and there is an edge between two nodes H if the corresponding spanning trees are neighbors.

Is it true that for any connected graph G , the resulting graph H is connected? Give a proof that H is always connected, or provide an example (with explanation) of a connected graph G for which H is not connected.

Answers: Yes, if a graph H contains multiple spanning trees, there must be at least 1 way to remove a unique edge e connecting nodes u and v and replacing it with another unique edge e' to connect u' to v , creating another connected graph. This would be by definition be 2 neighbors, which results in always having a connected graph H .